

What Do Wrist-Worn Stress Detectors Detect? Arousal Transfers Across Five Empatica E4 Datasets, Evaluative Stress Does Not

Alvaro Octavio Ibarra Reyes and Chen Pan

Abstract—Wrist-worn stress detectors are trained to separate stressor periods from rest and are reported as detecting stress. We ask what they detect. On five public Empatica E4 datasets with audited protocol-stage labels, leave-one-dataset-out evaluation on held-out subjects, per-subject normalisation and corrected significance tests, a model of heart-rate, heart-rate-variability and electrodermal features trained on four datasets loses only 0.015–0.05 balanced accuracy on the fifth; domain adaptation, few-shot personalisation and per-user thresholds add nothing significant. We score one fixed detector on non-stress arousal states it was never shown. Adding a state to the training negatives removes exercise false alarms (69% to 8%) but not false alarms on speech and arithmetic without evaluation (55% to 56%), hyperventilation (about 50%), manual tasks (43% to 33%) or film-induced emotions. At equal autonomic arousal the detector cannot separate stress from the same tasks without evaluation, hyperventilation, manual tasks or anger clips (AUROC 0.36–0.58, none different from chance after Holm correction), but separates it from exercise (0.89). Protocol position alone predicts stress better than physiology on the two hardest datasets, whose baselines are recorded while the sensor is still warming up. A test-side error budget attributes at most 0.06 of the 0.22–0.29 within-dataset error to recovery, sensor loss and onset latency; a strict responder filter recovers up to 0.09 on one dataset. Wrist detectors transfer because they detect autonomic arousal; the ceiling is stressor potency. We propose a specificity panel and a matched-arousal test as reporting standards and release audited labels, frozen subject splits and code.

Index Terms—Wearable stress detection, Empatica E4, autonomic arousal, cross-dataset generalization, specificity, electrodermal activity, heart rate variability.

I. INTRODUCTION

STRESS detection from wrist wearables is usually posed as a classification problem: windows recorded during a laboratory stressor are labelled stress, windows recorded during rest or a control task are labelled non-stress, and a model is trained to separate them. Reported accuracies on the standard datasets are high, and the models are described as detecting stress. Two observations motivate this paper. First, the cross-dataset literature disagrees about whether such models transfer: Prajod et al. [1] and Calza-Metre and Borzì [2] report large drops when the stressor changes, Xiao et al. [3] report balanced accuracy falling from 0.88 to 0.53–0.69 across seven wrist datasets, whereas Mishra et al. [4] found a cost of about 0.01 between two Empatica E4 studies with the same stressor, and Kwon et al. [5] found near-perfect transfer to

WESAD but near-chance results on Stress-Predict. Second, the wrist device measures pulse rate, a weak vagal proxy, wrist electrodermal activity (EDA), skin temperature and movement. None of these is stress; all of them respond to autonomic arousal from any cause. Whether a detector trained on stressor labels has learned anything beyond arousal has not, to our knowledge, been tested.

We address both observations with the same five public E4 datasets. After auditing every dataset’s labels against its descriptor paper, validating wrist cardiac features against chest electrocardiography (ECG), and fixing an evaluation protocol with subject-grouped splits and corrected tests, we show that the transfer cost is small once labels and per-subject normalisation are right, and that the remaining disagreement in the literature is largely about absolute difficulty rather than transfer. We then ask what transfers. A specificity panel scores one fixed detector on arousal states absent from training, and a matched-arousal test asks whether stress and non-stress arousal remain separable at equal arousal. Both answer that the detector separates exercise from stress, and fear clips once it has seen them, but not evaluative stress from the same tasks performed without evaluation, from hyperventilation, from manual tasks, or from anger clips. A time-only classifier, a warm-up analysis and a test-side error budget then locate the remaining within-dataset ceiling in the stressor rather than in the pipeline.

The contributions are: (1) a five-dataset leave-one-dataset-out benchmark with audited protocol-stage labels, showing a transfer cost of 0.015–0.05 balanced accuracy and diagnosing one apparent domain gap as a missing channel; (2) a specificity panel and a matched-arousal test showing that wrist detectors separate exercise but not non-evaluative arousal from stress; (3) a quantification of time in session and sensor warm-up as confounds, with a position-matched evaluation rule; (4) an error budget showing that recovery contamination, sensor loss and onset latency explain at most 0.06 of the within-dataset error, and that a strict responder filter recovers at most a further 0.09 on one dataset; (5) leakage-controlled negative results for domain adaptation, few-shot personalisation, per-user thresholds, self-rated stress intensity and a subject-level detectability trait. Labels, frozen splits and scripts are released. Section II reviews related work, Section III describes the datasets and the label audit, Section IV the features, normalisation and evaluation rules, Section V the results, and Section VI discusses implications, reporting norms and limitations.

The authors are with the Department of Electrical and Computer Engineering, The University of Texas at San Antonio, San Antonio, TX 78249 USA (e-mail: [email]@utsa.edu; [email]@utsa.edu).

II. RELATED WORK

A. Cross-dataset transfer

Kwon et al. [5] ran pooled leave-one-dataset-out (LODO) evaluation on three E4 corpora with logistic regression and gradient boosting and no per-subject normalisation, obtaining an area under the receiver operating characteristic curve (AUROC) of about 0.90 with WESAD as target but 0.52–0.57 on Stress-Predict and the Nurse dataset; their within-dataset AUROC on Stress-Predict was already 0.63, so the transfer cost there is 0.06–0.08. Calza-Metre and Borzì [2] report pairwise zero-shot transfer across four E4 corpora, with WESAD macro-F1 falling from 0.87 within to 0.63 from Campanella. Xiao et al. [3] train on one or two of seven wrist datasets and observe balanced accuracy 0.53–0.69 out of distribution. Mishra et al. [4], with per-participant normalisation and matched arithmetic stressors, found a median AUROC cost of about 0.01. Schreiber and Maleshkova [6] pool three datasets under leave-one-subject-out (LOSO) evaluation and name LODO as future work. Prajod et al. [1] attribute transfer failure to stressor type, on chest ECG. Our reading of this literature is that the reported failures come from low within-dataset ceilings, from stressor mismatch, or from the absence of per-subject normalisation, and the present study tests each.

B. Label quality and confounds

Several authors name label noise as a cause of poor transfer [5], [7] but none audits and fixes labels and measures the effect. Mishra et al. [4] exclude post-stressor rest because it carries residual arousal; Aydoğan and Villagra Povina [8] found independently that the PhysioNet stress sessions mix rest and stressor blocks. Richer et al. [9] note that comparing a stressor to a pre-stress baseline confounds the comparison with sequence effects. Tognotti et al. [10] show that test-inclusive baseline normalisation inflates balanced accuracy by 3–13 points, and Fecke and Rehof [11] normalise from a calibration baseline on chest ECG. Aydoğan and Villagra Povina [8] report that an XGBoost model labels 82.6% of PhysioNet exercise windows as stress, and 6–7% once exercise is a training class. Saeb et al. [12] and Bahameish et al. [13] document the inflation from record-wise cross-validation.

C. Personalisation and the arousal question

Adding a few labelled windows of a new user to a generic model is well established [14], [15], [16], [17]. Where calibration windows are taken chronologically, gains are small: Tervonen et al. [16] found retraining XGBoost gained 1.7 balanced-accuracy points on WESAD with 1–3 windows per class, and Akkaya [17] reports that the few-shot benefit did not survive a prospective chronological protocol. On wrist cardiac validity, Milstein and Gordon [18] found that the root mean square of successive differences (RMSSD) from the E4 correlated with ECG at $r = 0.42$ during speech, and Watanabe et al. [19] found the Trier Social Stress Test (TSST) to be the worst condition for E4 beat detection on WESAD. Sosa et al. [20], reviewing multimodal arousal classifiers, state that “physiological signals primarily index arousal and regulatory

TABLE I
HARMONISED CLASSES (SUBJECTS / 60-s WINDOWS AT 30-s STEP)

Dataset	Original label → harmonised class	Subj./win.
WESAD	TSST → Stress	15 / 311
	Baseline; Meditation; Amusement	15 / 566; 356; 165
PhysioNet	Stroop, TMCT, subtraction → Stress	35 / 232
	Baseline (3 min before first tag); First and second rest	35 / 174; 1193
	Aerobic; Anaerobic (whole sessions)	30 / 2143; 31 / 1620
Stress-Predict	Stroop; TSST interview → Stress	33 / 851
	Baseline; Relax → Rest	34 / 616; 1079
UBFC-Phys	Hyperventilation (probe only)	31 / 53
	T2, T3 test group → Stress	11 / 110
	T1 rest → Baseline	19 / 95
	T2, T3 control group → Control task (probe)	8 / 80
Campanella	Backward subtraction → Stress	29 / 102
	Baseline	29 / 116
	Lego tasks → Manual task (probe)	29 / 667
EPM-E4	Anger; Fear; Happiness; Sadness (probes)	33 / 33; 394; 297; 131

processes rather than affective valence” and that sympathetic activation “can therefore arise not only during stress, but also during non-stressful states”; the reviewed studies did not test classifiers on such states.

The gap is therefore twofold. No study has scored one fixed, subject-independent wrist detector across a panel of non-stress arousal states, or asked whether stress remains separable from those states at equal arousal; and no study has decomposed the within-dataset ceiling of wrist stress detection into recovery, sensor, onset, responder and stressor terms. The present paper addresses both on five public datasets with audited labels.

III. DATASETS AND LABEL AUDIT

Five public datasets recorded with the Empatica E4 (blood volume pulse (BVP) 64 Hz, EDA 4 Hz, skin temperature 4 Hz, accelerometer 32 Hz) contain both stress and non-stress periods: WESAD [21] (15 subjects; TSST speech and arithmetic before a panel, standing), the PhysioNet stress and exercise dataset of Hongn et al. [22] (35; Stroop, the Trier Mental Challenge Test (TMCT) of mental arithmetic under time pressure, opinion tasks, seated, plus aerobic and anaerobic sessions), Stress-Predict [23] (34 after exclusion; Stroop and a TSST-style interview, plus a hyperventilation block), UBFC-Phys [24] (19; speech and arithmetic, with a control group performing the same tasks without evaluation) and Campanella 2024 [25] (29; backward subtraction, plus Lego manual tasks with and without a countdown). EPM-E4 [26] (33 subjects with clip timing) supplies film-induced anger, fear, happiness and sadness and is used only as a probe. Table I gives the harmonised classes.

Every labelling decision was checked against the dataset’s descriptor paper. Seven corrections followed: in WESAD, labels had to be taken from the synchronised file rather than

mapped by length onto the raw E4 export, whose recording starts 16–26 min earlier; in PhysioNet, only the tagged stressor tasks count as stress (248 of 2,933 windows in the earlier whole-session labelling), one subject with a documented bad fit (f13) is excluded, and the 3 min before the first tag serve as baseline; in Stress-Predict, subject S01 violated the protocol and is dropped, and task boundaries whose tags are more than 180 s from the logged time are dropped; in Campanella, which logs no per-subject task markers, the subtraction span ends at a per-subject end estimated from the recording length. In addition, the Stress-Predict descriptor tags its 2-min hyperventilation block as a “stress-inducing task”; it is a respiratory manoeuvre with no evaluative component, and we keep it as a separate probe class. The three participants recorded in two files are treated as one subject each, so that no person appears on both sides of a split.

A cautionary note on evaluation. An earlier version of this pipeline used window-level train/test splitting with 50%-overlapping windows and reported 94% accuracy on a six-class task; the same model scores 49% (balanced accuracy 0.36) when 15% of subjects are held out (leakage_check.csv). Record-wise splitting [12] remains common in this field, and every number below uses subject-grouped splits.

IV. METHODS

A. Features

Windows are 60 s at a 30-s step, anchored inside one protocol stage. Cardiac features, the mean and standard deviation (SD) of heart rate (HR) and three heart-rate-variability (HRV) measures (RMSSD, the SD of normal-to-normal intervals (SDNN) and the proportion of successive differences above 50 ms (pNN50)), come from BVP beats detected with NeuroKit2 after band-pass filtering and artefact rejection; windows with fewer than 50% usable beats carry missing cardiac values. Against WESAD chest ECG on the same 1,398 windows, wrist heart rate has a mean absolute error of 2.4 bpm ($r = 0.94$); RMSSD is weak ($r = 0.53$, bias +34 ms), and only 37% of TSST windows keep usable photoplethysmography (PPG) because speech moves the wrist (heart-rate bias -7 bpm there). EDA features are mean, SD, range, slope, tonic and phasic means, and skin conductance response (SCR) count and amplitude. Temperature features are included only in the “physiology” set; the main set is HR/HRV/EDA, for reasons given in Sections V-B and V-E. Accelerometer features are used only in the modality comparison.

B. Normalisation

Features are z-scored per subject using all windows of that subject’s stress-protocol recording, without labels. This is transductive; two non-transductive variants are reported as sensitivity analyses: raw features, and causal scaling from strictly earlier windows only (the first four windows share the statistics of that opening period).

C. Task, models and evaluation

The main task is stress versus all non-stress states (baseline, rest, meditation, amusement). Baseline-versus-stress alone is confounded with time in session (Section V-E). The classifier is XGBoost with fixed hyperparameters; a nested-tuning comparison of XGBoost, random forest and a multilayer perceptron found one significant difference in 72 comparisons after Holm correction, on the nine-class affective task only and none on the stress tasks (significance_models.csv), so tuning is not repeated. Evaluation uses 20 subject-grouped splits with 20% of each dataset’s subjects held out, LOSO where per-subject scores are needed, and LODO in which the target’s held-out subjects are scored by a model trained (a) on the target’s other subjects, (b) on the other four datasets only, or (c) on both. Paired comparisons use Nadeau–Bengio corrected t -tests [27] with Holm correction; per-subject rates carry subject-bootstrap 95% confidence intervals (CI). The threshold is 0.5 throughout; balanced accuracy (BA) and AUROC are reported.

D. Four evaluation rules

Specificity panel. A fixed detector is scored on non-stress states absent from its task definition, reporting the false-stress rate ($p > 0.5$) under three training conditions: the state’s dataset absent from training; all datasets in training but the state itself absent; the state added to the training negatives. *Matched-arousal test.* A label-free arousal index (mean of per-subject z-scored heart rate and tonic EDA, or tonic EDA alone where heart rate is missing) is computed per window; stress windows are pooled with one non-stress arousal state, binned into arousal quintiles, and the AUROC of the stress score is computed within bins and averaged by bin size. An AUROC near 0.5 means the model has learned nothing beyond arousal magnitude for that pair. *Position-matched negatives.* Non-stress test windows are restricted to the 5th–95th percentile range of minutes-since-start spanned by that dataset’s stress windows. *Test-side error budget.* Filters applied to out-of-sample predictions without retraining, with gains Shapley-averaged over all filter orders.

V. RESULTS

A. Subject-independent performance

On the pooled baseline-versus-stress task of the three-dataset table (WESAD, PhysioNet, EPM-E4; 50 subjects) with raw (unnormalised) features and the fixed-hyperparameter model, 20 subject-grouped splits give BA 0.83 (95% CI 0.81–0.85) with all modalities, 0.81 (0.79–0.83) with physiology only and 0.75 (0.72–0.78) with the accelerometer only; LOSO agrees (0.85, 0.81, 0.75). No modality difference survives Holm correction ($p_{\text{holm}} = 1.0$, significance_modalities.csv), and a complete 108-configuration nested tuning run changes none of these conclusions. The model is therefore not detecting movement alone, and physiology carries most of the signal. On the main task and all five datasets, within-dataset BA with HR/HRV/EDA features is 0.916 (WESAD), 0.882 (Campanella), 0.877 (UBFC-Phys), 0.776 (PhysioNet) and 0.697

TABLE II
LEAVE-ONE-DATASET-OUT BALANCED ACCURACY ON HELD-OUT
SUBJECTS, STRESS VS ALL NON-STRESS

Target	HR/HRV/EDA			+ temperature		
	Within	Other 4	All 5	Within	Other 4	All 5
WESAD	0.916	0.884	0.921	0.900	0.840	0.876
PhysioNet	0.776	0.746	0.783	0.760	0.735	0.770
Stress-Predict	0.697	0.646	0.689	0.712	0.645	0.699
UBFC-Phys	0.877	0.837	0.839	0.877	0.589	0.867
Campanella	0.882	0.867	0.895	0.897	0.827	0.895




Fig. 1. Leave-one-dataset-out balanced accuracy on held-out subjects for the two feature sets. With temperature included, UBFC-Phys, which records no temperature, drops to 0.589 when trained on the other datasets although its AUROC stays at 0.931.

(Stress-Predict) (Table II). The two lowest ceilings belong to the two datasets whose stressors are seated, mostly silent and weakly evaluative, and whose physiological effect sizes are smallest: mean EDA rises by 0.4 z in PhysioNet against 2.2 z in WESAD, heart rate by 1.7 against 2.2 (stress_effect_size_by_dataset.csv).

B. Cross-dataset transfer

Table II and Fig. 1 give the LODO result. With HR/HRV/EDA features, a model trained on the other four datasets loses 0.03 (WESAD), 0.03 (PhysioNet), 0.05 (Stress-Predict), 0.04 (UBFC-Phys) and 0.015 (Campanella) BA against the within-dataset model on the same held-out subjects. Training on all five datasets matches or exceeds within-dataset training on WESAD, PhysioNet and Campanella and is within 0.04 on the other two (Stress-Predict 0.689 against 0.697, UBFC-Phys 0.839 against 0.877). These costs are within the error expected when each split holds out 3–7 subjects of a 15–35-subject dataset; we claim no large cost, not equality.

A missing channel posing as a domain gap. With temperature included, external UBFC-Phys falls to 0.589 BA while its AUROC stays 0.931 and an oracle threshold recovers 0.910 (threshold_transfer_probe_summary.csv). UBFC-Phys has no temperature in any window, so XGBoost’s missing-value branch pushes its mean stress score to 0.18 against 0.31–0.42 on the other targets. Ranking is intact; the threshold is not. Dropping temperature lifts external UBFC-Phys to 0.837. Elsewhere the remaining external loss is mostly ranking: an oracle threshold recovers only 0.02–0.05 with HR/HRV/EDA features, and a label-free cut matching the source stress rate hurts Campanella (0.867 to 0.732).

Normalisation. The small cost does not depend on transductive scaling. With raw features the external cost is −0.01 to 0.05 BA; causal scaling keeps it on WESAD (0.845 within,

TABLE III
FALSE-STRESS RATE ON STATES ABSENT FROM THE TASK DEFINITION
(HR/HRV/EDA; SUBJECT-BOOTSTRAP 95% CI)

Probed state (subjects / windows)	Dataset absent	Added as negative
PhysioNet aerobic / anaerobic (30–31 / 3763)	0.69 / 0.70	0.10 / 0.08
UBFC-Phys speech + arithmetic, no evaluation (8 / 80)	0.55 [0.46, 0.66]	0.56 [0.45, 0.67]
Stress-Predict hyperventilation (31 / 53)	0.42 [0.26, 0.59]	0.55 [0.36, 0.72]
Campanella Lego manual tasks (29 / 667)	0.43 [0.38, 0.48]	0.33 [0.27, 0.39]
EPM-E4 anger / fear / sadness / happiness (33)	0.42 / 0.35 / 0.37 / 0.24	0.35 / 0.25 / 0.30 / 0.18
Reference: in-task non-stress states	0.08–0.29	—
Reference: stress recall, by dataset	0.53–0.91	within −0.05

0.841 external), PhysioNet (0.774, 0.744) and Stress-Predict (0.688, 0.665) but loses about 0.10 externally on UBFC-Phys and Campanella, whose recordings are short (7–11 windows per subject). Whole-session scaling mainly raises the absolute level on both sides. No external difference between variants survives Holm correction across 60 tests.

Label fixes. Every external score rose after the seven corrections (for example UBFC-Phys 0.790 to 0.837, whose own data did not change), but on the two datasets where a paired test is possible only 1 of 16 differences survives Holm correction (WESAD external AUROC, +0.0075, $p_{\text{holm}} = 0.021$). In the two-dataset setting of Kwon et al. [5] (WESAD and Stress-Predict), the fixes change transfer by at most 0.02; our pipeline nevertheless reaches external AUROC 0.67 on Stress-Predict against their 0.56, and per-subject scaling raises WESAD as target from 0.834 to 0.914 AUROC. The label audit is a correctness contribution; the small transfer cost comes from features, task definition and normalisation.

C. Specificity panel

The LODO task keeps only baseline, rest, meditation and amusement as negatives; several recorded states were never scored. Table III and Fig. 2(a) give their false-stress rates. A model that never saw exercise calls 69–70% of PhysioNet exercise windows stress, matching the 71–83% of Aydoğan and Villagra Povina [8]; adding exercise to the training negatives collapses this to 8–10%, and the external PhysioNet drop is significant (BA 0.746 to 0.583, $p_{\text{holm}} = 0.011$). No other state behaves this way. The UBFC-Phys control group performs the same speech and arithmetic tasks as the test group without an evaluating panel and is called stress at 0.55 whether or not the model has seen it. Hyperventilation stays at about 0.5, Lego manual tasks fall only from 0.43 to 0.33, and the EPM-E4 emotion clips fall by 0.05–0.10. Adding a probed state to the negatives changes stress recall by at most −0.05 (UBFC-Phys

TABLE IV
MATCHED-AROUSAL TEST: AUROC OF THE STRESS SCORE INSIDE
AROUSAL QUINTILES

Stress versus	n	Pooled	Within-bin [CI]	p_{holm}	Disjoint
Lego manual task	667	0.754	0.355 [0.24, 0.54]	0.13	0.44
UBFC-Phys, no evaluation	80	0.587	0.577 [0.50, 0.67]	0.41	0.52
Hyperventilation	53	0.551	0.523 [0.44, 0.61]	0.62	0.52
Anger clips	33	0.643	0.390 [0.31, 0.52]	0.41	0.41
Fear clips	394	0.716	0.633 [0.58, 0.69]	0.004	0.65
Baseline and rest (reference)	3839	0.770	0.584 [0.54, 0.64]	0.004	0.57

Model trained with the probed state as a negative
Exercise 3763 0.947 **0.892** [0.84, 0.93] 0.003 0.88
Fear clips 394 0.804 **0.779** [0.74, 0.82] 0.003 0.78
Lego manual task 667 0.858 0.496 [0.44, 0.70] 0.94 0.52
UBFC-Phys, no evaluation 80 0.583 0.568 [0.46, 0.67] 0.64 0.50
Hyperventilation 53 0.568 0.560 [0.45, 0.67] 0.64 0.54
Anger clips 33 0.699 0.431 [0.33, 0.62] 0.64 0.52

Within-bin AUROC is weighted by bin size; CI is a subject bootstrap (2000 draws); p_{holm} is a permutation test against 0.5 (labels shuffled across subject-by-class cells within each bin, 2000 draws), Holm-corrected within each model type. “Disjoint” repeats the test with the two index features and `eda_mean` removed from the model.

0.74 to 0.70 with Lego added; PhysioNet 0.74 to 0.71 and Campanella 0.92 to 0.89 with the EPM-E4 clips added), and exercise negatives raise it (Stress-Predict 0.63 to 0.70, UBFC-Phys 0.74 to 0.79), so the failure is not a sensitivity trade-off. The hyperventilation (1.7 windows per subject) and UBFC-Phys control (8 subjects) cells are small; Lego and the fear clips carry the statistics.

D. Matched-arousal test

Table IV and Fig. 2(b) make the mechanism explicit. Once windows are binned by arousal, the external detector shows no separation of stress from a seated non-evaluative task beyond what arousal level provides: evaluated versus non-evaluated speech and arithmetic in UBFC-Phys gives a within-bin AUROC of 0.577, hyperventilation 0.523, Lego manual tasks 0.355 and anger clips 0.390, and none of these differs from chance after Holm correction (p_{holm} 0.13–0.62; `partA_inference.csv`). Adding the state to the training negatives does not change this (0.43–0.57, $p_{\text{holm}} \geq 0.64$). Two states are separable. Exercise reaches 0.892 within bins once the model has seen it ($p_{\text{holm}} = 0.003$), and fear clips reach 0.779; the bootstrap difference between the exercise AUROC and each non-separable pair excludes zero (lower bounds 0.18–0.36). For that model, repeating the test with a heart-rate-only index gives 0.91 and with an EDA-only index 0.82, so exercise stays separable whichever channel defines arousal; it is separated by its pattern across channels, not by arousal level. (For the external model, which has not seen exercise, the same indices give 0.65 and 0.54.)

Because the index shares two features with the model, the test was repeated with those features (and `eda_mean`) removed from the model (“Disjoint” in Table IV; `partA_disjoint.csv`). The conclusions are unchanged: UBFC-Phys control, hyperventilation, Lego and anger stay at 0.41–0.54, exercise at 0.88 and fear at 0.78. Two further variants, a heart-rate-only index with an EDA-only model and an EDA-only index with a cardiac-only model, give the

same pattern, except that Lego becomes separable (0.73–0.84) under the first variant, where the pair is restricted to the 259 stress and Lego windows with usable pulse (63% of Campanella windows have none) and is therefore not comparable. Exercise windows sit almost entirely above every stress window (median index 4.8 against 0.8), so its matched bins cover only the lowest three quintiles.

E. Time in session and sensor warm-up

Baseline is recorded first in every dataset, so any baseline-referenced comparison is confounded with time in session. Fig. 3(a) quantifies this: minutes since recording start, used as the only feature, reaches BA 0.83 on PhysioNet (physiology 0.78; better in 16 of 20 splits) and 0.95 on Stress-Predict (physiology 0.70; 20 of 20), ties on WESAD (0.90 against 0.92) and is perfect on Campanella. Given minutes as an extra feature, the within-dataset model exploits it (Stress-Predict 0.70 to 0.95) and then collapses externally (0.66 to 0.56; WESAD 0.90 to 0.45), because protocol timing does not transfer. Any within-dataset pipeline that leaks time, through raw temperature, drift features or the absence of per-subject normalisation, inherits this. The physiology model survives position-matched negatives with changes of at most -0.03 within and -0.04 externally (`time_probe_summary.csv`), so the LODO result is not riding on time; the matched negatives are mostly post-stressor rest, so this also bounds recovery contamination. A baseline-referenced feature set, by contrast, calls 50–83% of later non-stress windows stress across the WESAD–PhysioNet pair (`order_confound_check.csv`).

Time is informative on these datasets because baseline coincides with the warm-up phase of the sensor. Measured from the start of each E4 recording (`warmup_since_donning.csv`), subject-centred skin temperature rises at 0.02–0.08 °C/min in the first 10 min on all four datasets with a continuous clock keeps rising past 20 min on WESAD and rises until about minute 20 on Campanella before levelling off (+1.8 and +0.9 °C between minutes 1–5 and 30–40; WESAD subjects gain +0.3 to +6.5 °C over 30 min), and tonic EDA rises in the first 10 min on all four. Baseline falls inside that phase on Campanella (100% of baseline windows within 10 min of recording start), Stress-Predict (97%) and PhysioNet (68% within 10 min, 90% within 15), and stress arrives at a median of 7–28 min. WESAD is the exception: its synchronised protocol starts 16–26 min after the raw E4 recording (median 21 min, found by matching the wrist EDA stream), so its baseline begins 17–49 min after the device was switched on (median 33 min) and no baseline window falls within 15 min of recording start. WESAD is also the one dataset on which the time-only classifier does not beat physiology. Whether recording start equals donning is not documented by any dataset; on WESAD, every subject’s first minute of temperature reads 4–6 °C above skin, which shows the device was placed within that minute. This is the quantitative form of the sequence effect described by Richer et al. [9] and the reason temperature features track time and are excluded from the main feature set.




Fig. 2. (a) Specificity panel: false-stress rate of the HR/HRV/EDA detector on states absent from its task definition, under three training conditions, with subject-bootstrap 95% CIs. (b) Matched-arousal test: AUROC of the stress score within arousal quintiles for stress versus each state (external model; for exercise, the model that has seen exercise). Only exercise, and fear clips once in training, remain separable at equal arousal.




Fig. 3. (a) Balanced accuracy of a classifier that sees only minutes since recording start, against the physiology model within and across datasets. (b), (c) Subject-centred skin temperature and tonic EDA against minutes since the first protocol window (2-min bins); the shaded band marks the first 10 min, in which 100% of Campanella, 97% of Stress-Predict and 68% of PhysioNet baseline windows fall when measured from recording start (90% within 15 min). WESAD's protocol clock starts 16–26 min after its recording; EPM-E4's starts at recording.

F. Error budget and the chest-ECG bound

Table V and Fig. 4 test four explanations for the low ceilings on PhysioNet and Stress-Predict as test-side filters on LOSO and external predictions: F1 drops rest windows within 5 min of a stressor offset (recovery contamination); F2 keeps windows with cardiac coverage of at least 0.5 (sensor loss); F3 drops the first 60 s of each stressor (onset latency); F4 keeps only subjects whose own stressor-minus-baseline heart rate or tonic EDA exceeds 0.5 baseline SD (non-response; label-using, diagnostic only). Together they buy 0.05–0.06 BA on PhysioNet and Stress-Predict (0.02 for

the external model on Stress-Predict), and only F1 and F3 on Stress-Predict and F3 external on PhysioNet have confidence intervals excluding zero. The residual error of 0.17–0.23 (0.33 external on Stress-Predict) is explained by none of them. The responder filter is the one term sensitive to its definition (`partB_responder_sensitivity.csv`). At 0.5 or 1.0 baseline SD on either channel the two rules select the same subjects (32 of 35 PhysioNet, 32 of 34 Stress-Predict) and F4 gains 0.01–0.02. Requiring 1.0 SD on both channels keeps 17 of 35 PhysioNet subjects and raises within-dataset BA by 0.094 [0.05, 0.15] (external 0.065 [0.01, 0.13]); keeping the

TABLE V
TEST-SIDE ERROR BUDGET: BALANCED-ACCURACY GAIN OF EACH
FILTER, SHAPLEY-AVERAGED OVER 24 ORDERS (HR/HRV/EDA)

Dataset, model	All	F1	F2	F3	F4	Resid.
PhysioNet, within	0.776	0.003	0.007	0.020	0.020	0.175
PhysioNet, external	0.746	0.015	0.006	0.036	0.008	0.190
Stress-Predict, within	0.710	0.019	0.011	0.018	0.012	0.229
Stress-Predict, external	0.653	0.015	-0.015	0.012	0.005	0.329
WESAD, external	0.887	0	0.045	0	0.002	0.066




Fig. 4. Error budget. Each bar is 1-BA of the out-of-sample predictions; coloured segments are the Shapley-averaged gains of the four test-side filters, grey is the residual after all four. Negative gains (Stress-Predict external, F2) are drawn as zero.

top half of subjects by pooled response gives 0.062 [0.01, 0.13]. On Stress-Predict the strict rules keep 12–17 of 34 subjects and gain 0.03 within (CI includes zero) and -0.02 to +0.03 externally. Under the strict rules non-response is the largest single term on PhysioNet, about half of the four-filter total, yet the residual stays at 0.10–0.14 on PhysioNet and 0.21–0.35 on Stress-Predict. Onset latency and recovery are real but small: within-dataset Stress-Predict recall is 0.47 in the first minute of a stressor against 0.77 after 3 min, PhysioNet 0.63 in the first minute against 0.77 in the third, and rest false alarms peak in the first minute after offset (PhysioNet within 0.26 against 0.06–0.09 later); too few windows sit there to move BA.

On WESAD the only sizeable term is the sensor (F2, +0.045 external). Replacing wrist cardiac features with chest-ECG features on the same windows raises BA from 0.916 to 0.956 (AUROC 0.977 to 0.993), and chest ECG alone reaches 0.918, equal to the whole wrist pipeline; the gain concentrates in low-coverage windows, where wrist heart rate is absent (0.886 to 0.962). EDA alone gives 0.891. Motion-robust wrist heart rate is therefore worth at most about +0.04 on WESAD, and about 0.04 of error remains with chest-grade cardiac data.

TABLE VI
LEAKAGE-CONTROLLED NEGATIVE RESULTS (BALANCED ACCURACY,
HR/HRV/EDA UNLESS STATED)

Question	Result
Unsupervised domain adaptation (correlation alignment, maximum mean discrepancy, domain-adversarial and subject-adversarial networks), WESAD $\leftrightarrow$ PhysioNet, physiology	0.52–0.70 vs source-only 0.54–0.70; none reliably better
Few-shot on unseen dataset, first-in-time windows, $k = 1$ –5 per class	-0.04 to +0.04 refit gain; none significant after Holm
Few-shot, random windows	only significant gains (WESAD $k = 5$, +0.03); not deployable
Per-user threshold from calibration windows	hurts PhysioNet (-0.08 , $p_{\text{holm}} < 0.05$ at $k = 3$)
Self-rated stress rise vs window recall (PhysioNet 54 units, WESAD 15)	PhysioNet $\rho = -0.08$ [-0.42, 0.21] within, 0.02 [-0.30, 0.31] external; WESAD -0.49 [-0.87, 0.17] within, -0.60 [-0.84, -0.08] external (15 units); mild tasks not missed more
Detectability as a subject trait: recall across stressors	Stress-Predict $\rho = -0.14$ [-0.51, 0.24]; PhysioNet 0.44 [-0.04, 0.78]
Recall vs baseline HR, RMSSD, EDA, PPG coverage	$ \rho \leq 0.35$ with CIs including 0, except Stress-Predict external recall vs baseline EDA ($\rho = 0.48$ [0.18, 0.70])
Worst 20% of subjects' share of misses	37–61%
Exercise absent from training	69.5% of exercise windows called stress; 4–6% with exercise in training

G. Negative results

Table VI collects results that close directions the literature has proposed. Unsupervised domain adaptation on top of per-subject normalisation does not beat source-only training in either direction between WESAD and PhysioNet, consistent with Kwon et al. [5]. Few-shot personalisation of a model trained on the other four datasets, with each user's first k windows per class as calibration and a 30-s gap before any test window, adds at most 0.02–0.04 BA and nothing significant; drawing calibration windows at random from the recording produces the only significant gains, in line with Tervonen et al. [16] and Akkaya [17]. A per-user threshold hurts. Self-rated stress intensity does not predict which windows are missed: in PhysioNet, Stroop (mean rise 0.7 of 10) is detected as well as arithmetic (2.5); WESAD's negative correlation rests on the two subjects with the largest rise. The data give no support to detectability as a stable subject trait: recall is not correlated across stressors, is not predicted by baseline physiology or PPG coverage (one exception, Stress-Predict external recall against baseline EDA), and misses are only moderately concentrated; PhysioNet's within and external models do agree on who is hard ($\rho = 0.64$ [0.37, 0.82]).

VI. DISCUSSION

Arousal transfers, stress does not. The results fit one reading. A wrist device sees sympathetic arousal through a low-fidelity window. Once labels are protocol-anchored and features are normalised per subject, the marginal distributions of

that arousal are similar across laboratory datasets, so a detector transfers at small cost and domain adaptation has nothing to align. What the detector has learned is arousal magnitude, plus a cross-channel pattern that distinguishes exercise once exercise is in training; it has not learned evaluative threat, because at equal arousal it cannot tell an evaluated speech from an unevaluated one, a subtraction task from a Lego task, or a stressor from hyperventilation. Fear clips are the second separable state, and the disjoint variants place that residual in the cardiac features as much as in EDA, which is consistent with a distinct cardiac response to threat films rather than with an evaluative signature. The datasets with low ceilings are the ones whose stressors produce a fifth of the electrodermal response of the TSST; filters for recovery, sensor loss and onset recover at most 0.06 of that gap, and a strict responder filter at most 0.09 on PhysioNet and nothing on Stress-Predict. The mild-stress, few-shot and domain-adaptation nulls follow: self-report is the wrong ruler for an arousal detector, and personalisation cannot recover a response that did not occur.

Reporting norms. We propose that a wearable stress detector be published with a specificity panel: its false-stress rate under exercise, speech without evaluation, cognitive load without threat, and strong emotion, the way an assay reports cross-reactivity; and with a matched-arousal test where such states are available. We also propose three protocol rules: a settling period before baseline, or time since donning as a covariate; position-matched negatives, or at least a report of the time-only classifier, whenever baseline precedes the stressor; and exercise among the training negatives, or an activity gate, whenever the deployment includes movement.

What a decisive study needs. No public E4 dataset separates speech from evaluation within subject or counterbalances order. The between-subject UBFC-Phys contrast is the closest, and it is small. A within-subject design with speech and evaluation crossed, cycling at matched heart rate, cognitive load without feedback, Latin-square order with baselines at both ends, chest ECG and a palmar EDA reference would settle whether wrist EDA carries an evaluative-threat signature that the present data cannot resolve, and would provide the first order-matched baseline on the wrist.

Limitations. Whole-session per-subject normalisation is transductive; raw and causal variants keep the small transfer cost on long recordings but lose about 0.10 on short ones. The hyperventilation and UBFC-Phys control cells are small, and the UBFC-Phys contrast is between subjects. Wrist EDA agrees with palmar EDA at about $r = 0.3$ [18], so a null on the evaluative-threat signature may be instrument failure rather than absence of the signature. The arousal index shares two features with the model; the feature-disjoint rerun holds, but under the heart-rate-only variant it cuts the usable Lego and hyperventilation pairs by 40–66%. Responder filters use test labels and are diagnostic only; the strict definitions discard half the subjects. All data are laboratory recordings; ambulatory validity is untested. The pooled subject count (132 in the five stress datasets) is small for a field-level claim, and per-target costs of 0.015–0.05 are within the sampling error of 3–7 held-out subjects per split.

VII. CONCLUSION

Across five public E4 datasets with audited labels, a wrist stress detector built on HR, HRV and EDA features transfers to an unseen dataset at a cost of 0.015–0.05 BA, and none of the methods tested here, from domain adaptation to few-shot calibration, improved on that. What transfers is autonomic arousal. The detector separates exercise from stress once shown exercise, and fear clips once shown them, but at equal arousal it cannot separate evaluative stress from the same tasks without evaluation, from hyperventilation, from manual work or from anger clips. Protocol position predicts stress better than physiology on the hardest datasets, whose baselines are recorded during sensor warm-up, and the within-dataset ceiling appears to be set by the stressor’s potency and by subjects who do not respond, rather than by the pipeline. The next gains in wrist stress detection are more likely to come from protocol design, contextual sensing and honest abstention than from classifiers. Two studies follow from this work: a counterbalanced within-subject protocol that crosses speech with evaluation, adds matched-heart-rate exercise and non-evaluative cognitive load, and records chest ECG and palmar EDA alongside the wrist, to test whether wrist EDA carries an evaluative-threat signature at all; and an ambulatory validation in which the specificity panel is scored against labelled daily-life arousal, which no laboratory dataset can provide.

DATA AND CODE AVAILABILITY

The five datasets are available from their original sources under their own licences. The audited label tables, the 20 frozen subject-grouped splits, the feature-extraction and evaluation scripts, and the scripts for the specificity panel, matched-arousal test, time-in-session probe and error budget are released at <https://github.com/ALVAROI12/smartwatch-stress-detection> (branch `jbhi-revision`).

REFERENCES

- [1] P. Prajod, B. Mahesh, and E. André, “Stressor type matters!—Exploring factors influencing cross-dataset generalizability of physiological stress detection,” in *Proc. Int. Conf. Multimodal Interaction (ICMI)*, 2024, pp. 508–517, doi: 10.1145/3678957.3685741.
- [2] M. Calza-Metre and L. Borzi, “Machine learning-based automatic stress detection: Performance and generalization across datasets,” *Smart Health*, vol. 41, Art. no. 100693, 2026, doi: 10.1016/j.smhl.2026.100693.
- [3] Y. Xiao, H. Sharma, S. Kaur, D. Bergen-Cico, and A. Salekin, “Human heterogeneity invariant stress sensing,” *Proc. ACM Interact. Mob. Wearable Ubiquitous Technol.*, vol. 9, no. 3, Art. no. 3749465, pp. 1–42, 2025, doi: 10.1145/3749465.
- [4] V. Mishra, S. Sen, G. Chen, T. Hao, J. Rogers, C.-H. Chen, and D. Kotz, “Evaluating the reproducibility of physiological stress detection models,” *Proc. ACM Interact. Mob. Wearable Ubiquitous Technol.*, vol. 4, no. 4, Art. no. 147, 2020, doi: 10.1145/3432220.
- [5] N. Kwon, S. Yoon, J. Hur, and H. S. Kang, “When wellness stress detection enters healthcare workflows: An iso-heart-rate, cross-corpus evaluation of wrist wearables,” *Healthcare*, vol. 14, no. 15, Art. no. 2434, 2026, doi: 10.3390/healthcare14152434.
- [6] P. Schreiber and M. Maleshkova, “AutoStress benchmark: Evaluating factors that influence cross dataset generalizability in stress recognition,” in *Proc. IEEE Conf. Artif. Intell. (CAI)*, 2026, pp. 1335–1341, doi: 10.1109/CAI68641.2026.11536323.
- [7] P. Singh, A. Gupta, S. Jalan, M. Kumar, and P. Singh, “FEEL: Quantifying heterogeneity in physiological signals for generalizable emotion recognition,” arXiv preprint arXiv:2604.05926, 2026, doi: 10.48550/arXiv.2604.05926.

- [8] Y. Aydoğan and F. Villagra Povina, "Stress or arousal? Exercise confounding in wearable stress detection," *Med. Eng. Phys.*, vol. 147, Art. no. 095021, 2026, doi: 10.1088/1873-4030/aea41e.
- [9] R. Richer *et al.*, "Machine learning-based detection of acute psychosocial stress from body posture and movements," *Sci. Rep.*, vol. 14, Art. no. 8251, 2024, doi: 10.1038/s41598-024-59043-1.
- [10] A. Tognotti, C. F. Ottestanu, L. Anceschi, and C. Menon, "Baseline normalization choices inflate classification performance in wearable health monitoring: Quantification and mitigation strategies," *Front. Digit. Health*, vol. 8, Art. no. 1827279, 2026, doi: 10.3389/fdgth.2026.1827279.
- [11] M. Fecke and J. Rehof, "Machine learning-based stress detection using robust HRV features and person-specific normalization," *IEEE Access*, vol. 14, pp. 74450–74471, 2026, doi: 10.1109/ACCESS.2026.3692752.
- [12] S. Saeb, L. Lonini, A. Jayaraman, D. C. Mohr, and K. P. Kording, "The need to approximate the use-case in clinical machine learning," *GigaScience*, vol. 6, no. 5, 2017, doi: 10.1093/gigascience/gix019.
- [13] M. Bahameish, T. Stockman, and J. Requena Carrión, "Strategies for reliable stress recognition: A machine learning approach using heart rate variability features," *Sensors*, vol. 24, no. 10, Art. no. 3210, 2024, doi: 10.3390/s24103210.
- [14] C. L. Stewart, A. Folarin, and R. Dobson, "Personalized acute stress classification from physiological signals with neural processes," arXiv preprint arXiv:2002.04176, 2020, doi: 10.48550/arXiv.2002.04176.
- [15] Y. Han, P. Zhang, M. Park, and U. Lee, "Systematic evaluation of personalized deep learning models for affect recognition," *Proc. ACM Interact. Mob. Wearable Ubiquitous Technol.*, vol. 8, no. 4, Art. no. 206, pp. 1–35, 2024, doi: 10.1145/3699724.
- [16] J. Tervonen, E. Amparore, M. Botta, J. Närväinen, K. Pettersson, and J. Mäntyjärvi, "Data-constrained personalization of cognitive state detection with feature-based and foundation models," *User Model. User-Adapt. Interact.*, 2026, doi: 10.1007/s11257-026-09444-w.
- [17] A. Akkaya, "Calibration, not architecture, limits cross-subject wearable stress detection: A multi-dataset, decision-focused evaluation with label-efficient recalibration," *BMC Med. Inform. Decis. Mak.*, 2026, doi: 10.1186/s12911-026-03842-1 (abstract only).
- [18] N. Milstein and I. Gordon, "Validating measures of electrodermal activity and heart rate variability derived from the Empatica E4 utilized in research settings that involve interactive dyadic states," *Front. Behav. Neurosci.*, vol. 14, Art. no. 148, 2020, doi: 10.3389/fnbeh.2020.00148.
- [19] Y. Watanabe, N. Yamane, A. Sathyanarayana, V. Mishra, and M. S. Goodwin, "Beyond motion artifacts: Optimizing PPG preprocessing for accurate pulse rate variability estimation," in *Companion Proc. ACM Int. Joint Conf. Pervasive Ubiquitous Comput. (UbiComp Companion)*, 2025, pp. 1176–1182, doi: 10.1145/3714394.3756241.
- [20] S. Sosa, A. K. Fontecchio, E. G. Chrysikou, and J. S. Atchison, "Beyond EDA: A systematic review of multimodal sympathetic nervous system arousal classification for stress detection," *Sensors*, vol. 26, no. 5, Art. no. 1584, 2026, doi: 10.3390/s26051584.
- [21] P. Schmidt, A. Reiss, R. Dürichen, C. Marberger, and K. Van Laerhoven, "Introducing WESAD, a multimodal dataset for wearable stress and affect detection," in *Proc. Int. Conf. Multimodal Interaction (ICMI)*, 2018, doi: 10.1145/3242969.3242985.
- [22] A. Hongn, F. Bosch, L. E. Prado, J. M. Ferrández, and M. P. Bonomini, "Wearable physiological signals under acute stress and exercise conditions," *Sci. Data*, vol. 12, Art. no. 520, 2025, doi: 10.1038/s41597-025-04845-9.
- [23] T. Iqbal *et al.*, "Stress monitoring using wearable sensors: A pilot study and Stress-Predict dataset," *Sensors*, vol. 22, no. 21, Art. no. 8135, 2022, doi: 10.3390/s22218135.
- [24] R. Meziati Sabour, Y. Benezeth, P. De Oliveira, J. Chappé, and F. Yang, "UBFC-Phys: A multimodal database for psychophysiological studies of social stress," *IEEE Trans. Affect. Comput.*, vol. 14, no. 1, pp. 622–636, 2023, doi: 10.1109/TAFFC.2021.3056960.
- [25] S. Campanella, A. Altaieb, A. Belli, P. Pierleoni, and L. Palma, "PPG and EDA dataset collected with Empatica E4 for stress assessment," *Data Brief*, vol. 53, Art. no. 110102, 2024, doi: 10.1016/j.dib.2024.110102.
- [26] F. M. Garcia-Moreno and M. Badenes-Sastre, "EmoPulse Moments E4 dataset (EPM-E4): An exhaustive collection of emotion-related data from Empatica E4 wearable," Zenodo, 2020, doi: 10.5281/zenodo.8431638.
- [27] C. Nadeau and Y. Bengio, "Inference for the generalization error," *Mach. Learn.*, vol. 52, no. 3, pp. 239–281, 2003, doi: 10.1023/A:1024068626366.